



Faculty of Engineering
Cairo University

Machine learning Project

Binary Prediction of Poisonous Mushrooms



Brief Problem Description

Building machine learning model to predict whether a mushroom is edible or poisonous based on its physical characteristics.

Problem Motivation

Accurately determining whether a mushroom is edible or poisonous is critical for both foragers and food safety professionals. Misidentification can lead to severe health consequences, including poisoning and even death. Traditional identification methods rely on expert knowledge, which is not always accessible to the general public.

Dataset

<https://www.kaggle.com/competitions/playground-series-s4e8/data>

Evaluation Metrics

Accuracy, Precision, Recall, F1 Score

Matthews correlation coefficient (MCC): is a metric used to evaluate the quality of binary classifications, especially when the classes are imbalanced.

$$MCC = \frac{TP \cdot TN - FP \cdot FN}{\sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}}$$

Project Pipeline

Data loading

Train.csv file shape: (3116945, 22)

Test.csv file shape: (2077964, 21)

Data Splitting

Split train to train and validation (80 , 20) with stratification on the target feature ("class")

Data cleaning

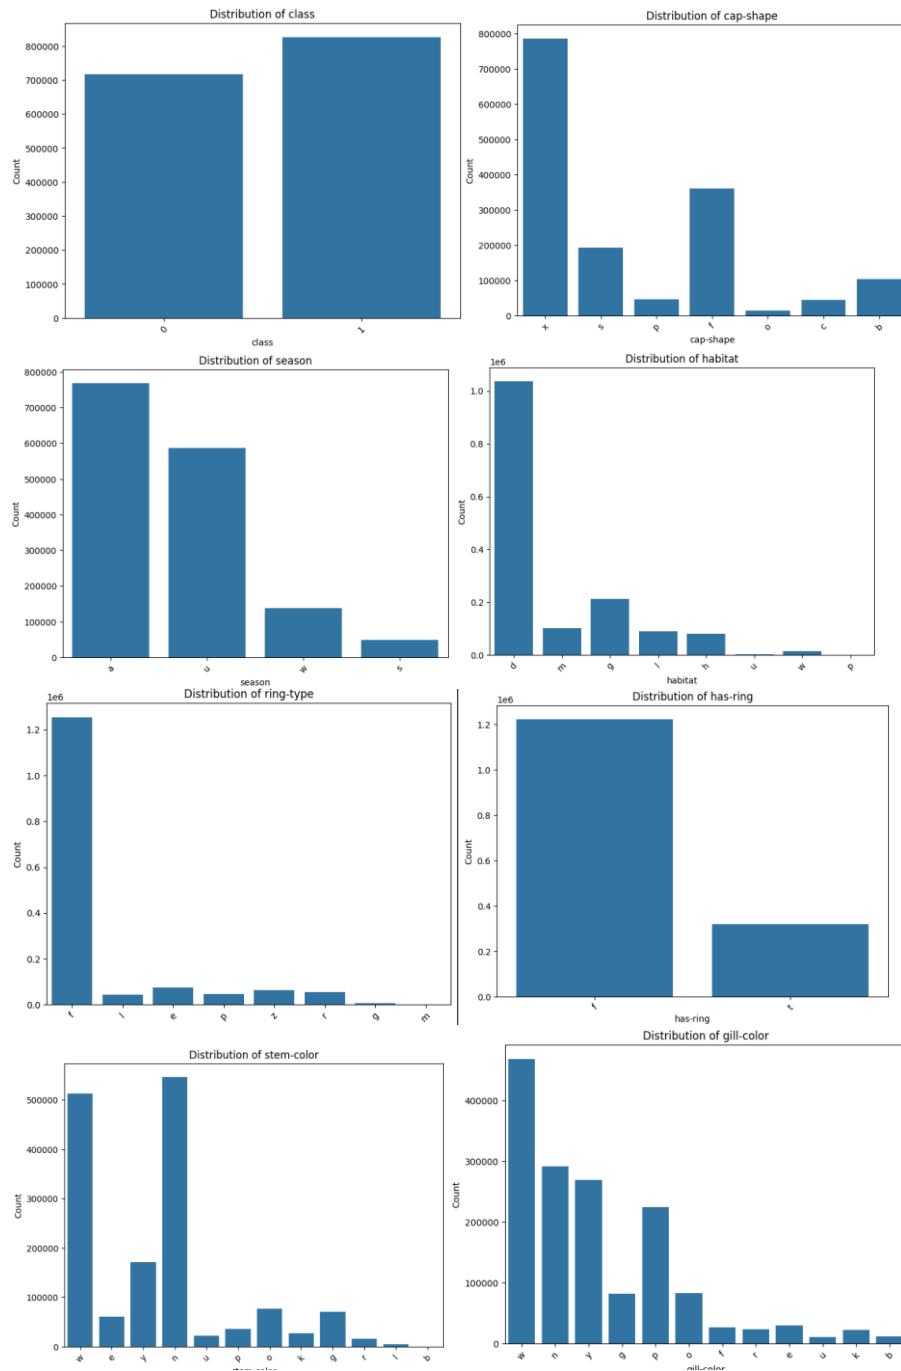
1. Dropping features with null values > 40% -> 6 features been dropped
2. The data had some noise in the categorical features values -> we took the noise values and combine them as 'other' values and then to NaN -> so we can fill them with meaningful values
3. Taking version of train dataset and dropping rows with null so we can train our target encoder -> we used target encoder because the values of the categorical features was nominal and we have more than 7 unique values on average for each categorical feature so we cannot use either OHE or Label-Encoder

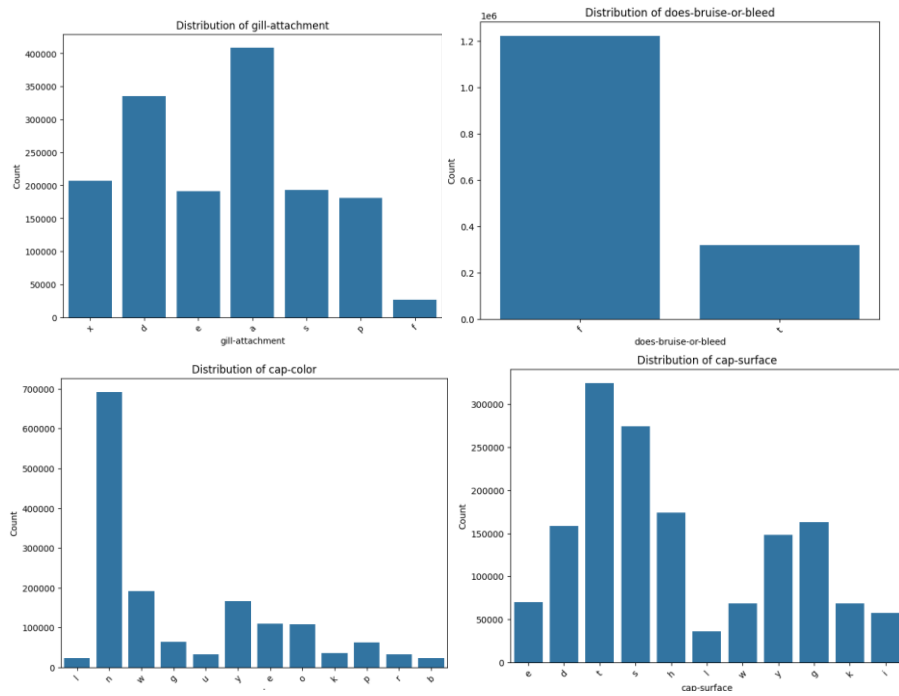
4. Transform the categorical features of train , test and val with the target encoder
5. Impute the missing values (and noise) with Simple imputer using mean (failed to use KNN imputer took a lot of time and Ram crash)
6. Now we have cleaned data

Data visualization

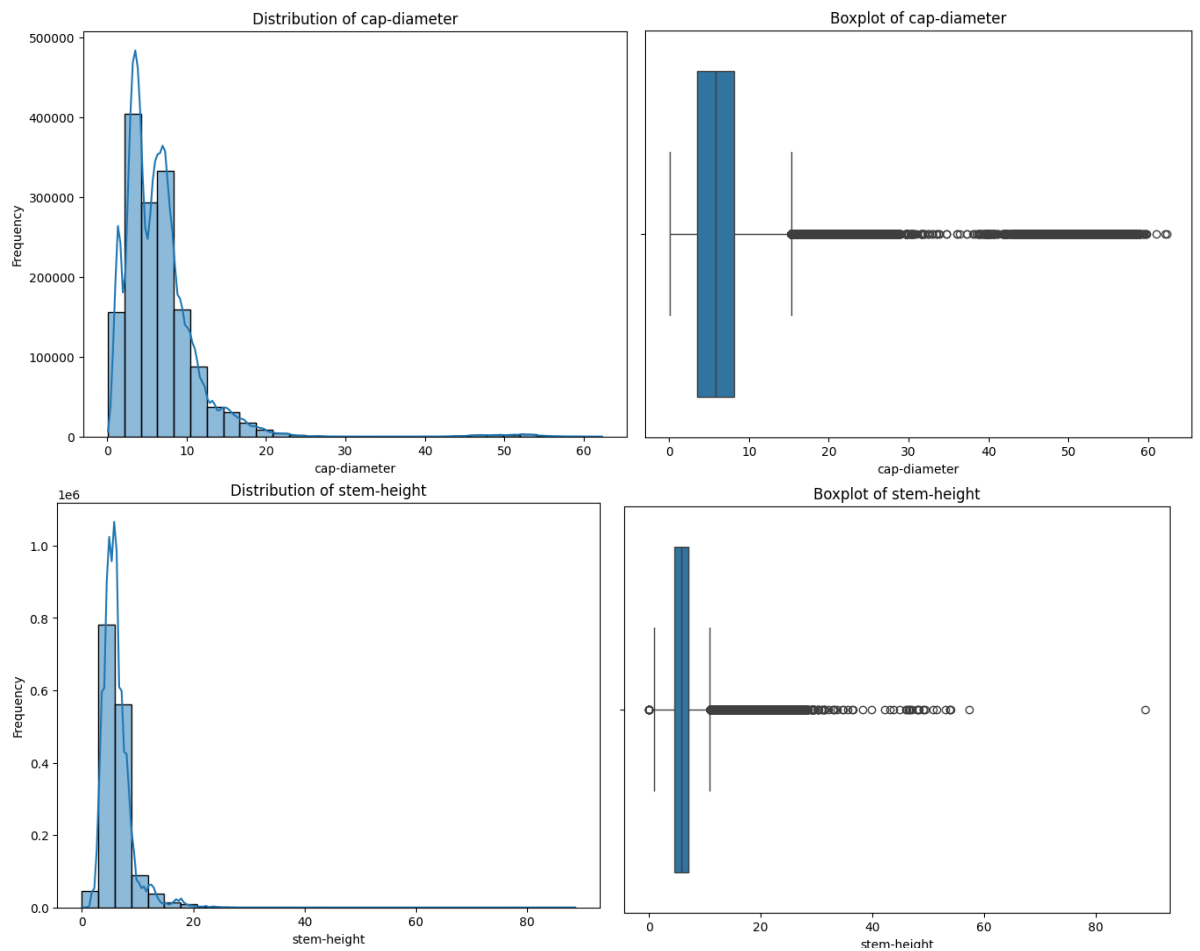
1. Univariate Analysis

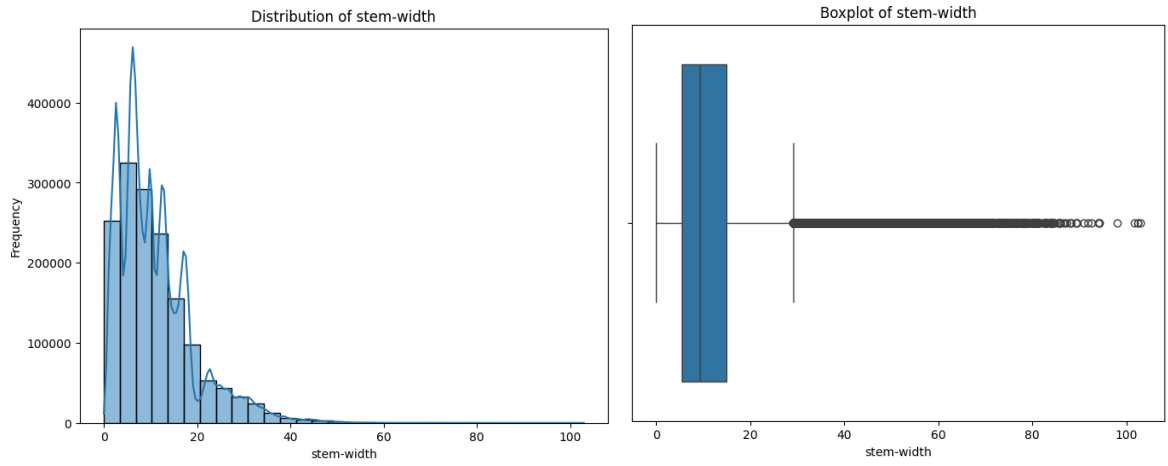
- Categorical features





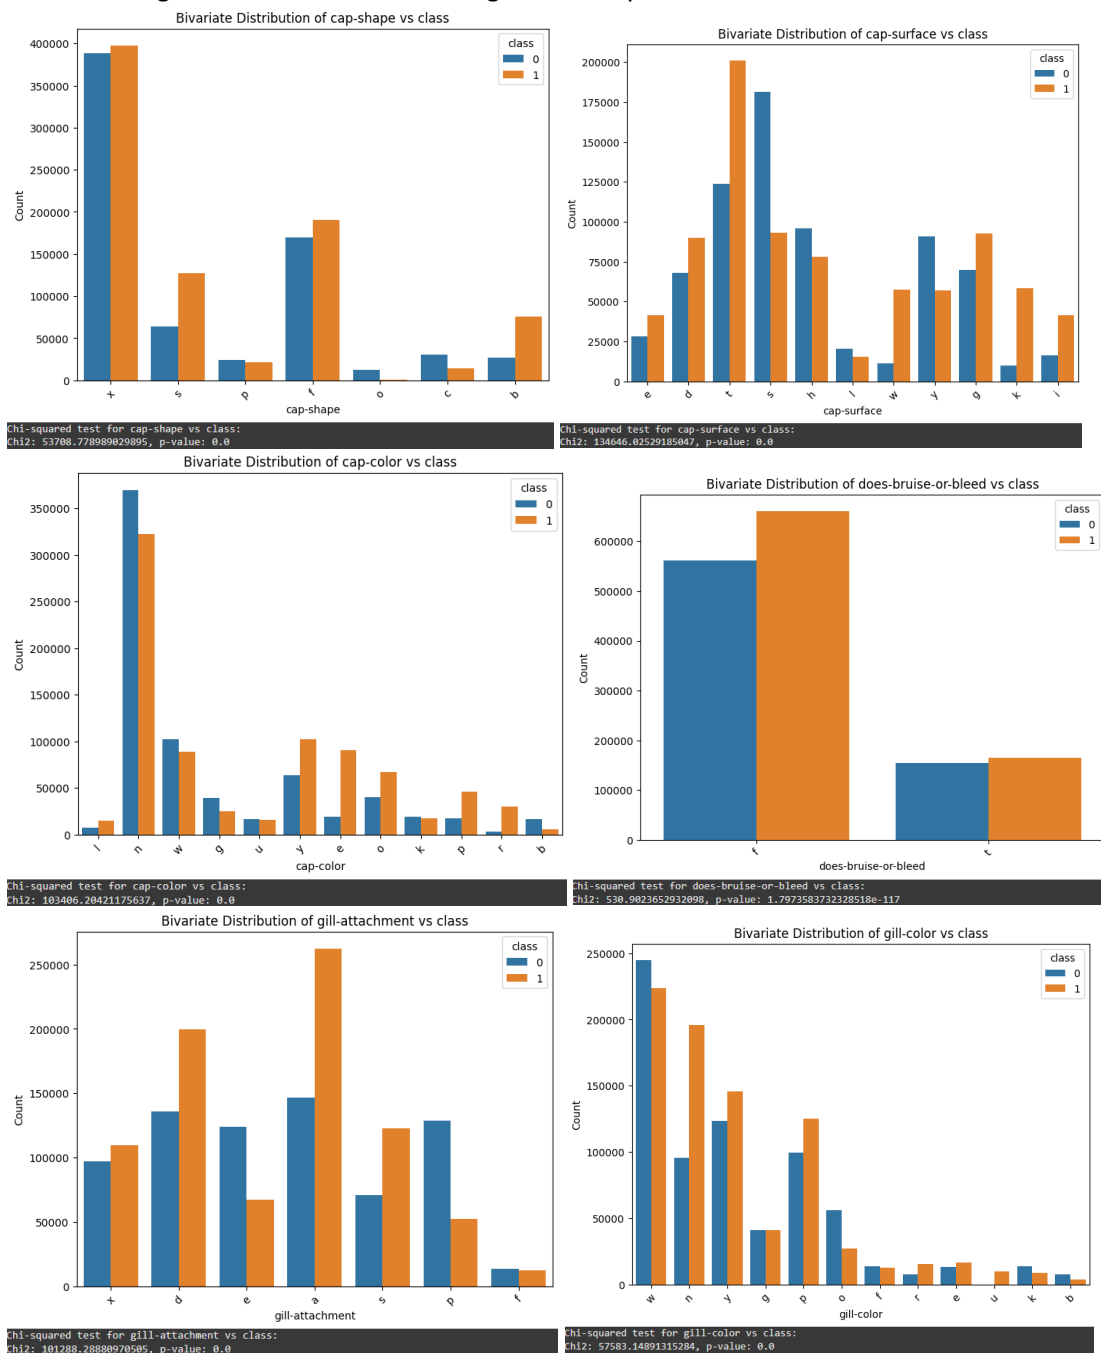
- Numerical features

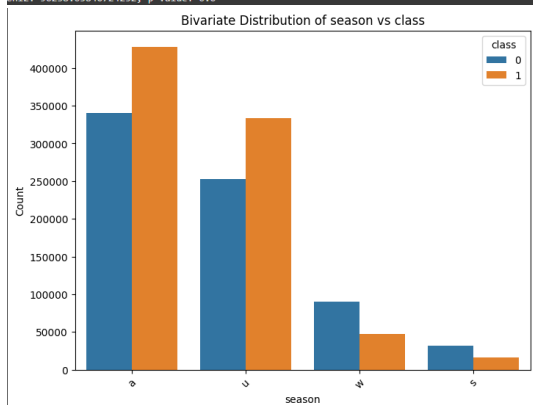
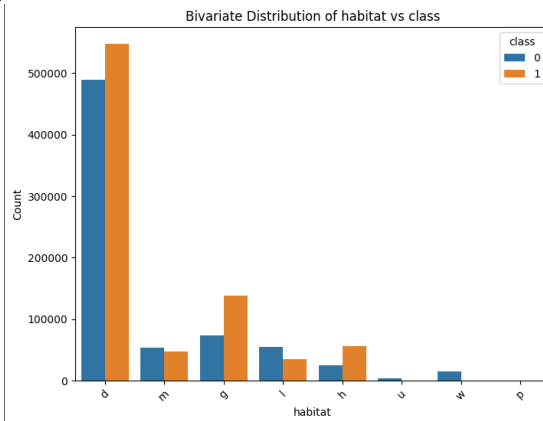
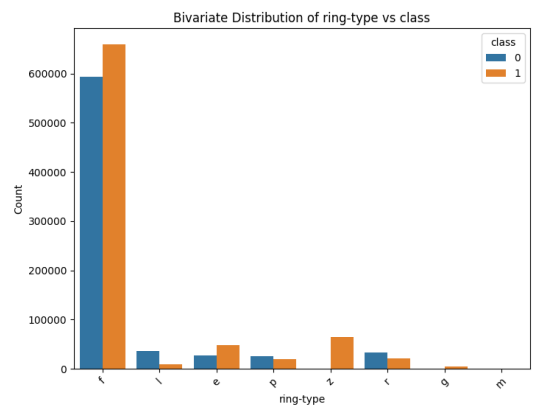
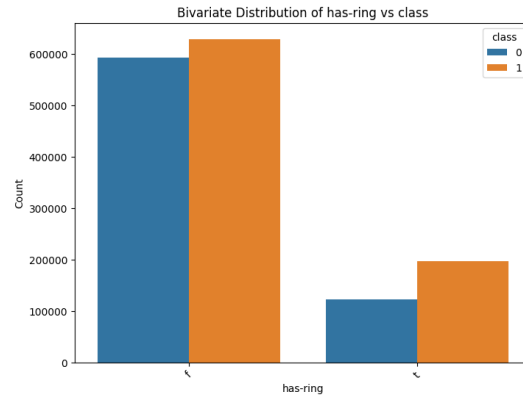
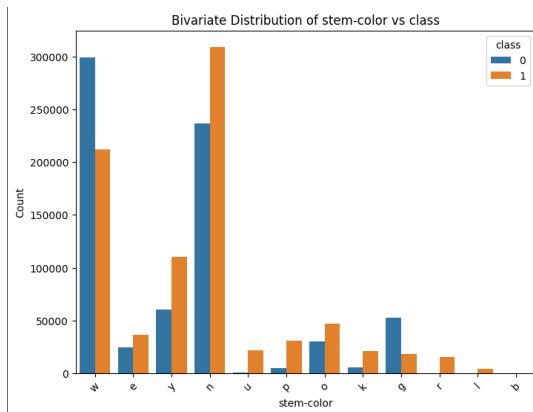




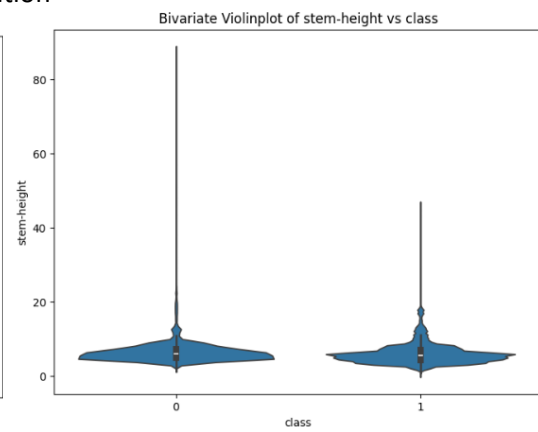
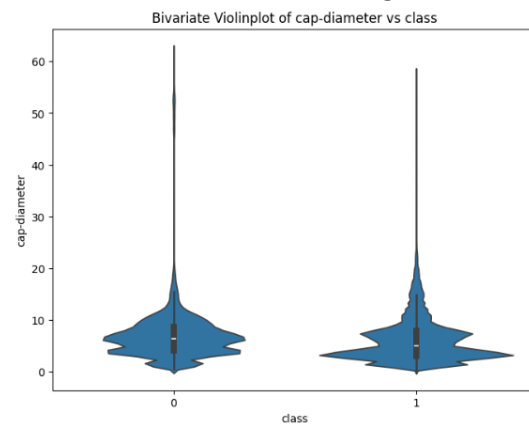
2. Bivariant Analysis


- Categorical features with the target & Chi-Squared test



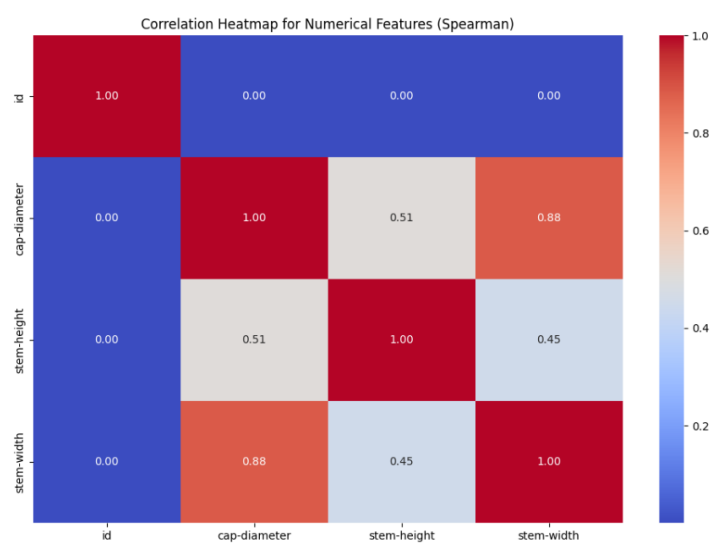
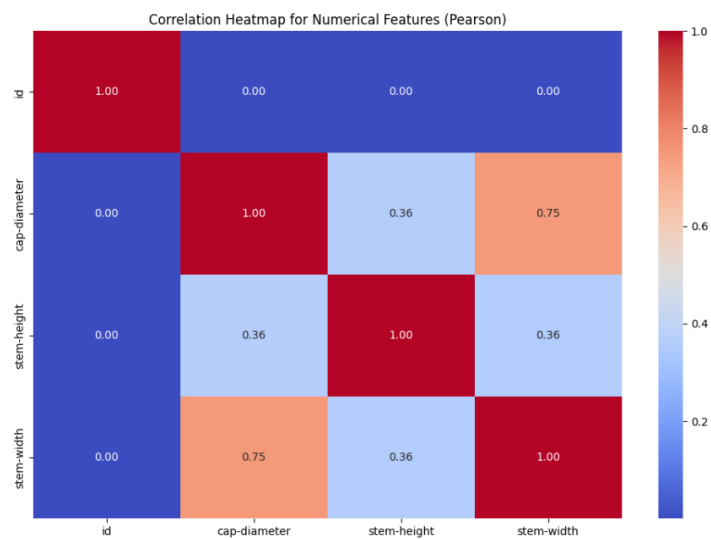


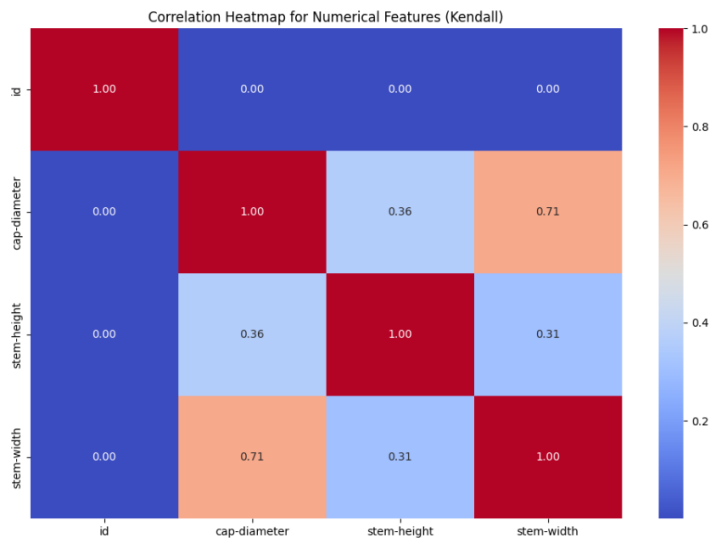
- Numerical features with the target & Correlation





- Multicollinearity





Feature selection

From the visuals we have only one problem

stem-width & cap-diameter is high correlated -> will drop one and keep the highest correlated with the target

```
Pearson correlation between 'stem-width' and 'class': -0.16934260357984848
Spearman correlation between 'stem-width' and 'class': -0.22160116090702292
Kendall correlation between 'stem-width' and 'class': -0.18099223947462714

Pearson correlation between 'cap-diameter' and 'class': -0.1622962730474319
Spearman correlation between 'cap-diameter' and 'class': -0.19821371816220826
Kendall correlation between 'cap-diameter' and 'class': -0.16192168554973005
```

We dropped cap-diameter

Outliers removing

We remove the outliers using IQR method

Modeling

logReg

We apply Standard Scaler first on the data

Parm : (random_state=42, max_iter=1000 ,n_jobs=-1 ,C=1,class_weight=None)


```

=== Logistic Regression ===
---- Train Metrics ----
[[754808 316765]
 [319804 943523]]
      precision    recall  f1-score   support

     0       0.70      0.70      0.70    1071573
     1       0.75      0.75      0.75    1263327

   accuracy          0.73          0.73          0.73    2334900
  macro avg          0.73          0.73          0.73    2334900
weighted avg          0.73          0.73          0.73    2334900

Train Accuracy: 0.7274
Train MCC: 0.4512
---- Validation Metrics ----
[[202443  79867]
 [ 98546 242533]]
      precision    recall  f1-score   support

     0       0.67      0.72      0.69    282310
     1       0.75      0.71      0.73    341079

   accuracy          0.71          0.71          0.71    623389
  macro avg          0.71          0.71          0.71    623389
weighted avg          0.72          0.71          0.71    623389

Validation Accuracy: 0.7138
Validation MCC: 0.4265

```

Note: high bias & small (Acceptable) variance

Param: (random_state=42, max_iter=1000 ,n_jobs=-1,solver='lbfgs',C=10,class_weight=None)

```

=== Logistic Regression (C=10 , new solver algorithm) ===
---- Train Metrics ----
[[754808 316765]
 [319804 943523]]
      precision    recall  f1-score   support

     0       0.70      0.70      0.70    1071573
     1       0.75      0.75      0.75    1263327

   accuracy          0.73          0.73          0.73    2334900
  macro avg          0.73          0.73          0.73    2334900
weighted avg          0.73          0.73          0.73    2334900

Train Accuracy: 0.7274
Train MCC: 0.4512
---- Validation Metrics ----
[[202442  79868]
 [ 98546 242533]]
      precision    recall  f1-score   support

     0       0.67      0.72      0.69    282310
     1       0.75      0.71      0.73    341079

   accuracy          0.71          0.71          0.71    623389
  macro avg          0.71          0.71          0.71    623389
weighted avg          0.72          0.71          0.71    623389

Validation Accuracy: 0.7138
Validation MCC: 0.4265

```

Note: still high bias & small (Acceptable) variance

Action: try polynomial features to increase the model complexity

```

poly_logreg_pipeline = Pipeline([
    ("poly", PolynomialFeatures(degree=2, interaction_only=False, include_bias=False)),
    ("scaler", StandardScaler()),
    ("logreg", LogisticRegression(max_iter=1000, random_state=42))
])

```

```

=== Logistic Regression with Polynomial Features ===
---- Train Metrics ----
[[ 846059 225514]
 [ 202853 1060474]]
      precision    recall  f1-score   support

     0       0.81      0.79      0.80    1071573
     1       0.82      0.84      0.83    1263327

 accuracy          0.82    2334900
 macro avg          0.82    2334900
weighted avg          0.82    2334900

Train Accuracy: 0.8165
Train MCC: 0.6301
---- Validation Metrics ----
[[214695 67615]
 [ 55411 285668]]
      precision    recall  f1-score   support

     0       0.79      0.76      0.78    282310
     1       0.81      0.84      0.82    341079

 accuracy          0.80    623389
 macro avg          0.80    623389
weighted avg          0.80    623389

Validation Accuracy: 0.8026
Validation MCC: 0.6007

```

Note: improved but still high bias

Action: increase the polynomial degree -> Data is huge and takes a lot of time

SVM

Parm: SVC(kernel='linear')

Takes forever

Action: perform on sample of the data

Parm: SVC(kernel='linear') with 10k row only

```

=== SVM (Balanced Sample) ===
---- Train Metrics ----
[[3370 1219]
 [1400 4011]]
      precision    recall  f1-score   support

     0       0.71      0.73      0.72     4589
     1       0.77      0.74      0.75     5411

 accuracy          0.74    10000
 macro avg          0.74    10000
weighted avg          0.74    10000

Train Accuracy: 0.7381
Train MCC: 0.4745
---- Validation Metrics ----
[[210982 71328]
 [104306 236773]]
      precision    recall  f1-score   support

     0       0.67      0.75      0.71    282310
     1       0.77      0.69      0.73    341079

 accuracy          0.72    623389
 macro avg          0.72    623389
weighted avg          0.72    623389

Validation Accuracy: 0.7183
Validation MCC: 0.4396

```

Note: high Bias , (Acceptable) Variance

Action : more complex kernel

Parm: SVC(kernel='rbf')

```

=== SVM (Balanced Sample (rbf) ) ===
---- Train Metrics ----
[[4256 333]
 [ 279 5132]]
      precision    recall  f1-score   support

     0       0.94      0.93      0.93      4589
     1       0.94      0.95      0.94      5411

   accuracy          0.94      10000
  macro avg       0.94      0.94      0.94      10000
 weighted avg     0.94      0.94      0.94      10000

Train Accuracy: 0.9388
Train MCC: 0.8767
---- Validation Metrics ----
[[250070 32240]
 [ 17811 323268]]
      precision    recall  f1-score   support

     0       0.93      0.89      0.91     282310
     1       0.91      0.95      0.93     341079

   accuracy          0.92      623389
  macro avg       0.92      0.92      0.92     623389
 weighted avg     0.92      0.92      0.92     623389

Validation Accuracy: 0.9197
Validation MCC: 0.8382

```

Note: improvement on bias and (Acceptable) variance

Action: increase the size of the sample 5 times

```

=== SVM (Balanced Sample (rbf & increased samples)) ===
---- Train Metrics ----
[[21712 1235]
 [ 1002 26051]]
      precision    recall  f1-score   support

     0       0.96      0.95      0.95     22947
     1       0.95      0.96      0.96     27053

   accuracy          0.96      50000
  macro avg       0.96      0.95      0.95      50000
 weighted avg     0.96      0.96      0.96      50000

Train Accuracy: 0.9553
Train MCC: 0.9099
---- Validation Metrics ----
[[259352 22958]
 [ 12530 328549]]
      precision    recall  f1-score   support

     0       0.95      0.92      0.94     282310
     1       0.93      0.96      0.95     341079

   accuracy          0.94      623389
  macro avg       0.94      0.94      0.94     623389
 weighted avg     0.94      0.94      0.94     623389

Validation Accuracy: 0.9431
Validation MCC: 0.8853

```

Note: improvement on bias and variance but takes long time

Action : increase the sample size more -> **Huge amount of time**

Action : try to use tree based model

Decision Tree

Parm: DecisionTreeClassifier(random_state=42)

```

=== Decision Tree ===
---- Train Metrics ----
[[1071450 123]
 [ 780 1262547]]
      precision    recall  f1-score   support

     0       1.00      1.00      1.00    1071573
     1       1.00      1.00      1.00    1263327

 accuracy          1.00
 macro avg          1.00
 weighted avg       1.00

Train Accuracy: 0.9996
Train MCC: 0.9992
---- Validation Metrics ----
[[275625 6685]
 [ 6571 334508]]
      precision    recall  f1-score   support

     0       0.98      0.98      0.98    282310
     1       0.98      0.98      0.98    341079

 accuracy          0.98
 macro avg          0.98
 weighted avg       0.98

Validation Accuracy: 0.9787
Validation MCC: 0.9571

```

Note : gives good results -> low bias , high variance (overfitting)

Action: try to reduce overfitting -> (tune max_depth , ccp_alpha)

Parm: DecisionTreeClassifier(random_state=42,max_depth=14)

```

=== Decision Tree ===
---- Train Metrics ----
[[1057976 13597]
 [ 20048 1243279]]
      precision    recall  f1-score   support

     0       0.98      0.99      0.98    1071573
     1       0.99      0.98      0.99    1263327

 accuracy          0.99
 macro avg          0.99
 weighted avg       0.99

Train Accuracy: 0.9856
Train MCC: 0.9710
---- Validation Metrics ----
[[277913 4397]
 [ 6176 334903]]
      precision    recall  f1-score   support

     0       0.98      0.98      0.98    282310
     1       0.99      0.98      0.98    341079

 accuracy          0.98
 macro avg          0.98
 weighted avg       0.98

Validation Accuracy: 0.9830
Validation MCC: 0.9658

```

Note : reduce the overfitting

Parm: DecisionTreeClassifier(random_state=42,ccp_alpha=0.001)

ccp_alpha: Controls complexity through cost-complexity pruning-> Higher values prune more of the tree.

```

=== Decision Tree ===
---- Train Metrics ----
[[1021272  50301]
 [ 56021 1207306]]
      precision    recall  f1-score   support

         0         0.95      0.95      0.95    1071573
         1         0.96      0.96      0.96    1263327

 accuracy          0.95          0.95          0.95    2334900
 macro avg          0.95          0.95          0.95    2334900
 weighted avg       0.95          0.95          0.95    2334900

Train Accuracy: 0.9545
Train MCC: 0.9084
---- Validation Metrics ----
[[267647  14663]
 [ 18967 322112]]
      precision    recall  f1-score   support

         0         0.93      0.95      0.94    282310
         1         0.96      0.94      0.95    341079

 accuracy          0.95          0.95          0.95    623389
 macro avg          0.95          0.95          0.95    623389
 weighted avg       0.95          0.95          0.95    623389

Validation Accuracy: 0.9461
Validation MCC: 0.8914

```

Note: reduce overfitting but increase bias

Parm: DecisionTreeClassifier(random_state=42, ccp_alpha=0.0001)

```

=== Decision Tree ===
---- Train Metrics ----
[[1045527  26046]
 [ 25901 1237426]]
      precision    recall  f1-score   support

         0         0.98      0.98      0.98    1071573
         1         0.98      0.98      0.98    1263327

 accuracy          0.98          0.98          0.98    2334900
 macro avg          0.98          0.98          0.98    2334900
 weighted avg       0.98          0.98          0.98    2334900

Train Accuracy: 0.9778
Train MCC: 0.9552
---- Validation Metrics ----
[[274568   7742]
 [ 7783 333296]]
      precision    recall  f1-score   support

         0         0.97      0.97      0.97    282310
         1         0.98      0.98      0.98    341079

 accuracy          0.98          0.98          0.98    623389
 macro avg          0.97          0.97          0.97    623389
 weighted avg       0.98          0.98          0.98    623389

Validation Accuracy: 0.9751
Validation MCC: 0.9497

```

Note: reduce overfitting , (Acceptable) bias -> but we can increase bias more

Action : more complex -> Random forest

Random Forest

Parm: RandomForestClassifier(random_state=42) -> default parm

```

=== Random Forest ===
---- Train Metrics ----
[[1071149  424]
 [ 518 1262809]]
      precision    recall  f1-score   support

     0       1.00      1.00      1.00    1071573
     1       1.00      1.00      1.00    1263327

 accuracy          1.00      1.00      1.00    2334900
 macro avg          1.00      1.00      1.00    2334900
 weighted avg       1.00      1.00      1.00    2334900

Train Accuracy: 0.9996
Train MCC: 0.9992
---- Validation Metrics ----
[[278978  3332]
 [ 3561 337518]]
      precision    recall  f1-score   support

     0       0.99      0.99      0.99    282310
     1       0.99      0.99      0.99    341079

 accuracy          0.99      0.99      0.99    623389
 macro avg          0.99      0.99      0.99    623389
 weighted avg       0.99      0.99      0.99    623389

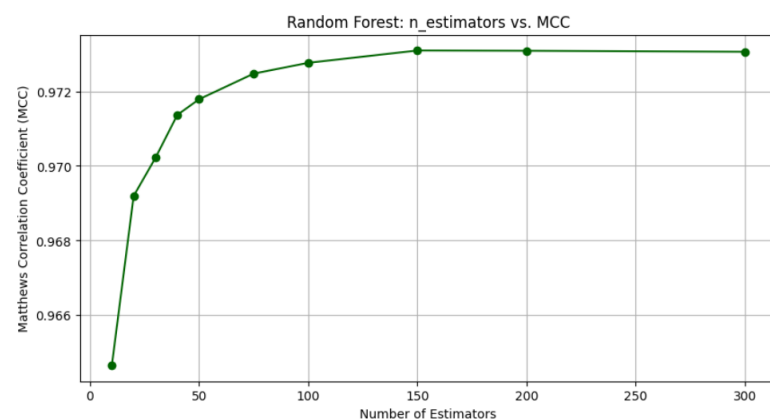
Validation Accuracy: 0.9889
Validation MCC: 0.9777

```

Note: low bias & small (Acceptable) variance

Action : increase the number of estimators (study the best number of estimators first)

Working on sample of the data 200K row



Parm: RandomForestClassifier(n_estimators=150, random_state=42)

```

=== Random Forest (double estimators) ===
---- Train Metrics ----
[[1071163  410]
 [ 494 1262833]]
      precision    recall  f1-score   support

     0       1.00      1.00      1.00    1071573
     1       1.00      1.00      1.00    1263327

 accuracy          1.00      1.00      1.00    2334900
 macro avg          1.00      1.00      1.00    2334900
 weighted avg       1.00      1.00      1.00    2334900

Train Accuracy: 0.9996
Train MCC: 0.9992
---- Validation Metrics ----
[[278996  3314]
 [ 3508 337571]]
      precision    recall  f1-score   support

     0       0.99      0.99      0.99    282310
     1       0.99      0.99      0.99    341079

 accuracy          0.99      0.99      0.99    623389
 macro avg          0.99      0.99      0.99    623389
 weighted avg       0.99      0.99      0.99    623389

Validation Accuracy: 0.9891
Validation MCC: 0.9779

```

Note: small change

Action : try Ensemble Models

Bagging

Parm : BaggingClassifier(estimator =DecisionTreeClassifier(), n_estimators=50, random_state=42,n_jobs=-1)

```
=== Bagging (Decision Tree) ===
---- Train Metrics ----
[[1071056  517]
 [ 729 1262598]]
      precision    recall  f1-score   support

     0       1.00      1.00      1.00    1071573
     1       1.00      1.00      1.00    1263327

 accuracy          1.00      1.00      1.00    2334900
 macro avg          1.00      1.00      1.00    2334900
 weighted avg       1.00      1.00      1.00    2334900

Train Accuracy: 0.9995
Train MCC: 0.9989
---- Validation Metrics ----
[[278523  3787]
 [ 4254 336825]]
      precision    recall  f1-score   support

     0       0.98      0.99      0.99    282310
     1       0.99      0.99      0.99    341079

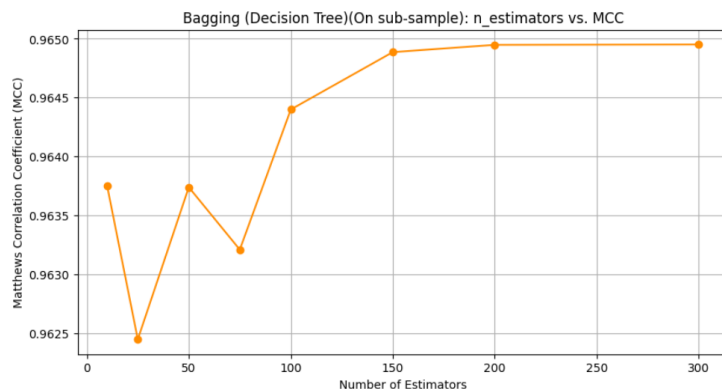
 accuracy          0.99      0.99      0.99    623389
 macro avg          0.99      0.99      0.99    623389
 weighted avg       0.99      0.99      0.99    623389

Validation Accuracy: 0.9871
Validation MCC: 0.9740
```

Note : converges really well with small number of estimators relative to the other models

Action : increase the number of estimator (study the best number first)

On sample of the data only 200K row



Parm: BaggingClassifier(estimator =DecisionTreeClassifier(), n_estimators=200, random_state=42,n_jobs=-1)

```
=== Bagging (Decision Tree)(200 estimators) ===
---- Train Metrics ----
[[1071157  416]
 [ 488 1262839]]
      precision    recall  f1-score   support

     0       1.00      1.00      1.00    1071573
     1       1.00      1.00      1.00    1263327

 accuracy          1.00      1.00      1.00    2334900
 macro avg          1.00      1.00      1.00    2334900
 weighted avg       1.00      1.00      1.00    2334900

Train Accuracy: 0.9996
Train MCC: 0.9992
---- Validation Metrics ----
[[278501  3809]
 [ 4074 337005]]
      precision    recall  f1-score   support

     0       0.99      0.99      0.99    282310
     1       0.99      0.99      0.99    341079

 accuracy          0.99      0.99      0.99    623389
 macro avg          0.99      0.99      0.99    623389
 weighted avg       0.99      0.99      0.99    623389

Validation Accuracy: 0.9874
Validation MCC: 0.9745
```

Note : small change

Adaboost

Parm: default parm

Note: bad performance

Action: increase the number of estimators

Parm: AdaBoostClassifier(n_estimators=100,random_state=42)

```
=== AdaBoost ===
---- Train Metrics ----
[[ 817145  254428]
 [ 258688 1004639]]
      precision    recall  f1-score   support

     0       0.76      0.76      0.76    1071573
     1       0.80      0.80      0.80    1263327

 accuracy          0.78    2334900
 macro avg       0.78      0.78      0.78    2334900
 weighted avg    0.78      0.78      0.78    2334900

Train Accuracy: 0.7802
Train MCC: 0.5576
---- Validation Metrics ----
[[215320  66990]
 [ 72270 268809]]
      precision    recall  f1-score   support

     0       0.75      0.76      0.76    282310
     1       0.80      0.79      0.79    341079

 accuracy          0.78    623389
 macro avg       0.77      0.78      0.77    623389
 weighted avg    0.78      0.78      0.78    623389

Validation Accuracy: 0.7766
Validation MCC: 0.5500
```

Note: high bias

Action : increase the number of estimator more -> **huge amount of time needed**

XGBoost

Parm: XGBClassifier(n_estimators=100, learning_rate=0.1, use_label_encoder=False, eval_metric='logloss', random_state=42)

```
=== XGBoost ===
---- Train Metrics ----
[[1049511  22062]
 [ 30124 1233203]]
      precision    recall  f1-score   support

     0       0.97      0.98      0.98    1071573
     1       0.98      0.98      0.98    1263327

 accuracy          0.98    2334900
 macro avg       0.98      0.98      0.98    2334900
 weighted avg    0.98      0.98      0.98    2334900

Train Accuracy: 0.9776
Train MCC: 0.9550
---- Validation Metrics ----
[[274710  7600]
 [ 8112 332967]]
      precision    recall  f1-score   support

     0       0.97      0.97      0.97    282310
     1       0.98      0.98      0.98    341079

 accuracy          0.97    623389
 macro avg       0.97      0.97      0.97    623389
 weighted avg    0.97      0.97      0.97    623389

Validation Accuracy: 0.9748
Validation MCC: 0.9491
```

Note: very low variance , needs to decrease bias && it was fast

Action : increase the number of estimator directly

Parm: xgb_model = XGBClassifier(objective='binary:logistic', use_label_encoder=False, n_estimators=500)


```

=== XGBoost ===
---- Train Metrics ----
[[1061342  10231]
 [ 13669 1249658]]
      precision    recall  f1-score   support

     0       0.99      0.99      0.99    1071573
     1       0.99      0.99      0.99    1263327

 accuracy
macro avg      0.99      0.99      0.99    2334900
weighted avg    0.99      0.99      0.99    2334900

Train Accuracy: 0.9898
Train MCC: 0.9794
---- Validation Metrics ----
[[279049  3261]
 [ 3719 337360]]
      precision    recall  f1-score   support

     0       0.99      0.99      0.99    282310
     1       0.99      0.99      0.99    341079

 accuracy
macro avg      0.99      0.99      0.99    623389
weighted avg    0.99      0.99      0.99    623389

Validation Accuracy: 0.9888
Validation MCC: 0.9774

```

Note: very low variance, very low bias

Parm: feval=mcc_eval, # custom MCC metric with 1000 estimator

```

=== XGBoost custom MCC ===
---- Train Metrics ----
[[1062312  9261]
 [ 12465 1250862]]
      precision    recall  f1-score   support

     0       0.99      0.99      0.99    1071573
     1       0.99      0.99      0.99    1263327

 accuracy
macro avg      0.99      0.99      0.99    2334900
weighted avg    0.99      0.99      0.99    2334900

Train Accuracy: 0.9907
Train MCC: 0.9813
---- Validation Metrics ----
[[279121  3189]
 [ 3501 337578]]
      precision    recall  f1-score   support

     0       0.99      0.99      0.99    282310
     1       0.99      0.99      0.99    341079

 accuracy
macro avg      0.99      0.99      0.99    623389
weighted avg    0.99      0.99      0.99    623389

Validation Accuracy: 0.9893
Validation MCC: 0.9783

```

Note: low variance, very low bias

Best model we can use is XGBoost -> let's tune some of his parameters using StratifiedKFold CV and GridsearchCV









```

xgb = XGBClassifier(
    objective='binary:logistic',
    use_label_encoder=False,
    eval_metric='logloss',
    verbosity=0,
    n_jobs=-1
)
param_grid = {
    'max_depth': [6, 10, 14],
    'min_child_weight': [5, 7],
    'gamma': [1e-6, 1e-4],
    'subsample': [0.7, 0.8],
    'colsample_bytree': [0.6],
    'reg_alpha': [0.1],
    'n_estimators': [50]
}
cv = StratifiedKFold(n_splits=3, shuffle=True, random_state=42)
mcc_scorer = make_scorer(matthews_corrcoef)
grid_search = GridSearchCV(
    estimator=xgb,
    param_grid=param_grid,
    scoring=mcc_scorer,
    n_jobs=-1,
    cv=cv,
    verbose=2
)

```

```
Best parameters: {'colsample_bytree': 0.6, 'gamma': 0.0001, 'max_depth': 14, 'min_child_weight': 5, 'n_estimators': 50, 'reg_alpha': 0.1, 'subsample': 0.8}
Best cross-validation score: 0.9797767095341166
```

Test Scores

 submission_CV_best_all_W_outliers.csv Complete (after deadline) · 6h ago	0.97998	0.98013	<input type="checkbox"/>
 submission_rf_150_estimator_all_W_outliers.csv Complete (after deadline) · 8h ago	0.97917	0.97902	<input type="checkbox"/>
 submission_rf_100_estimator_W_outliers.csv Complete (after deadline) · 20h ago	0.97892	0.97884	<input type="checkbox"/>
 submission.csv Complete (after deadline) · 2d ago · xgb_1000_estimator.csv XGboost with 80% of the train with no outliers	0.97863	0.97864	<input type="checkbox"/>
 submission_xgb_500_estimator_all_W_outliers.csv Complete (after deadline) · 9h ago	0.97853	0.97854	<input type="checkbox"/>
 submission_xgb_500_estimator_W_outliers.csv Complete (after deadline) · 9h ago	0.97843	0.97842	<input type="checkbox"/>
 submission_rf_200_estimator.csv Complete (after deadline) · 20h ago · random forest with 200 estimators	0.97779	0.97767	<input type="checkbox"/>
 submission_bagging_200_estimator.csv Complete (after deadline) · 9h ago	0.97454	0.97475	<input type="checkbox"/>

With less than 0.5% difference between the first score on the lead board (using AutoML)

Conclusion

- Tree based model was more powerful in our problem
- Xgboost is robust to overfitting relative to the models we used
- It always best practise to test the number of estimators relative to the evaluation metric to choose the best number of estimator that will lead to the best generalization
- Outliers can be not removed when you use tree based model and it can help the model performance
- Tree based model is much faster than logreg , svm model
- After getting the best model performance do intensive parm. tuning using gridSearchCV to get even more better result on your model
- Train the model on the whole dataset train , validation before submitting the test file on the competition it will give more better results
- Endcoding the features with the correct way (target encoder in our problem) can prevent you from getting really high dim data
- Splitting the data with stratification will help the model performance when dealing with imbalanced data